

main

June 16, 2023

1 Descripción del dataset

Este dataset concierne a las redes de metro globales. Se estudia los kilómetros, paradas, el uso, año de construcción y de última ampliación así como aspectos relacionados con la ciudad donde se encuentra: población, localización geográfica, país.

Tras limpiar estos datos se analizan aspectos como el número de viajes por habitante, latitudes con alta densidad de uso (este asiático), ciudades con mas estaciones, etc.

Las bases de datos originales están localizables en: - del metro <https://www.kaggle.com/datasets/drahulsingh/metro-systems-worldwide> - de ciudades <https://simplemaps.com/data/world-cities>

2 Integración y selección de los datos de interés a analizar

```
[1]: import numpy as np
import pandas as pd
from itertools import chain

import matplotlib.pyplot as plt
from matplotlib import ticker

plt.style.use('seaborn-v0_8-whitegrid')
import seaborn as sns

import warnings
import copy
```

2.1 Carga de bases de datos

Cargamos dataset del metro (de <https://www.kaggle.com/datasets/drahulsingh/metro-systems-worldwide>).

```
[2]: df_metro = pd.read_csv("../dataset/Metro-Systems-Worldwide.csv")
```

```
[3]: df_metro.rename(columns={'Country\r\nregion': 'Country',
                             'Service\r\nopened': 'Year opened',
                             'Last\r\nexpanded': 'Last year expanded',
```

```

        'Annual ridership\r\n(millions)': 'Annual ridership_
↪(millions)'
    }, inplace=True)

```

```
[4]: df_metro.head()
```

```

[4]:
      City      Country      Name  Year opened \
0   Algiers    Algeria  Algiers Metro    2011
1  Buenos Aires  Argentina  Buenos Aires Underground    1913
2   Yerevan    Armenia   Yerevan Metro    1981
3   Sydney    Australia   Sydney Metro    2019
4   Vienna     Austria   Vienna U-Bahn    1978

      Last year expanded  Stations      System length  Annual ridership (millions)
0              2018         19  18.5 km (11.5 mi)          45.3 (2019)
1              2019         78  56.7 km (35.2 mi)          74.0 (2020)
2              1996         10   12.1 km (7.5 mi)          23.3 (2022)
3                -         13    36 km (22 mi)          16.3 (2022)
4              2017         98  83.3 km (51.8 mi)         459.8 (2019)

```

Cargamos dataset de ciudades (de <https://simplemaps.com/data/world-cities>):

```

[5]: df_cities = pd.read_csv("../dataset/worldcities.csv")
df_cities =
↪df_cities[["city_ascii", "country", "admin_name", "population", "lat", "lng"]]
df_cities.head()

```

```

[5]:
  city_ascii  country  admin_name  population    lat    lng
0    Tokyo     Japan    Tōkyō    37732000.0  35.6897  139.6922
1   Jakarta  Indonesia    Jakarta    33756000.0  -6.1750  106.8275
2    Delhi     India     Delhi    32226000.0  28.6100   77.2300
3  Guangzhou    China  Guangdong    26940000.0  23.1300  113.2600
4    Mumbai     India  Mahārāshtra    24973000.0  19.0761   72.8775

```

2.2 Combinar ambos datasets

2.2.1 Campos del merge

Pero hay ciudades homónimas en muchos países (en ambos dataframes):

```
[6]: df_cities["city_ascii"].value_counts(), df_metro["City"].value_counts()
```

```

[6]: (Santa Cruz      17
      San Fernando   16
      Santa Ana      15
      Santa Maria    14
      San Juan       14
      ..
      Arsenyev       1

```

```

Panjacent      1
Kleve          1
Venkatagiri    1
Nordvik        1
Name: city_ascii, Length: 41140, dtype: int64,
Seoul          3
Tokyo          3
New York City  3
London         2
Manila         2
..
Medellín       1
Prague         1
Copenhagen     1
Santo Domingo  1
Hanoi          1
Name: City, Length: 194, dtype: int64)

```

Incluso mas complicado, hay ciudades repetidas en dataset del metro:

```

[7]: unique_cities_metro = df_metro.groupby(['City', 'Country']).agg({'Name': lambda x: x.index.tolist()}).reset_index()
unique_cities_metro.rename(columns={'Name': 'Original indexes'}, inplace=True)
unique_cities_metro['Count'] = [len(k) for k in unique_cities_metro['Original indexes']]
mask_cities_metro = unique_cities_metro['Count']>1
unique_cities_metro[mask_cities_metro]

```

```

[7]:
      City      Country Original indexes  Count
11  Bangkok      Thailand      [173, 174]      2
94   London  United Kingdom      [185, 186]      2
100  Manila    Philippines      [150, 151]      2
120  New York City  United States  [194, 195, 196]      3
130  Philadelphia  United States      [197, 198]      2
151   Seoul      South Korea  [136, 137, 138]      3
174   Tokyo        Japan    [124, 125, 126]      3
190  Yokohama        Japan    [127, 128]      2

```

```

[8]: # repeated cities vs sum
unique_cities_metro[mask_cities_metro]['Count'].shape[0] , \
unique_cities_metro[mask_cities_metro]['Count'].sum()

```

```

[8]: (8, 19)

```

Comprobamos que se debe que en una misma ciudad-país puede haber varios sistemas de metro:

```

[9]: indexes = list(chain(
    from_iterable(unique_cities_metro[mask_cities_metro]['Original indexes']))

```

```
df_metro.iloc[indexes ,:]
```

```
[9]:
```

	City	Country	Name \
173	Bangkok	Thailand	BTS Skytrain
174	Bangkok	Thailand	Metropolitan Rapid Transit
185	London	United Kingdom	London Underground
186	London	United Kingdom	Docklands Light Railway
150	Manila	Philippines	Manila Light Rail Transit System
151	Manila	Philippines	Manila Metro Rail Transit System
194	New York City	United States	New York City Subway
195	New York City	United States	Staten Island Railway
196	New York City	United States	PATH
197	Philadelphia	United States	SEPTA
198	Philadelphia	United States	PATCO Speedline
136	Seoul	South Korea	Seoul Metropolitan Subway
137	Seoul	South Korea	Korail metro lines
138	Seoul	South Korea	Shinbundang Line (Neo Trans)
124	Tokyo	Japan	Toei Subway
125	Tokyo	Japan	Tokyo Metro
126	Tokyo	Japan	Rinkai Line
127	Yokohama	Japan	Yokohama Municipal Subway
128	Yokohama	Japan	Minatomirai Line

	Year opened	Last year expanded	Stations	System length \
173	1999	2021	60	68.2 km (42.4 mi)
174	2004	2019	53	71 km (44 mi)
185	1863	2021	272	402 km (250 mi)
186	1987	2011	45	34 km (21 mi)
150	1984	2021	33	37.2 km (23.1 mi)
151	1999	2000	13	16.9 km (10.5 mi)
194	1904	2017	424	399 km (248 mi)
195	1925	2017	21	22.5 km (14.0 mi)
196	1908	1937	13	22.2 km (13.8 mi)
197	1907	1973	72	59.1 km (36.7 mi)
198	1936	1980	13	22.9 km (14.2 mi)
136	1974	2022	279	345.3 km (214.6 mi)
137	1994	2022	86	151.7 km (94.3 mi)
138	2011	2022	16	33.4 km (20.8 mi)
124	1960	2002	99	109.0 km (67.7 mi)
125	1927	2020	142	195.1 km (121.2 mi)
126	1996	2002	8	12.2 km (7.6 mi)
127	1972	2008	40	53.4 km (33.2 mi)
128	2004	2008	6	4.1 km (2.5 mi)

	Annual ridership (millions)
173	236.9 (2020)
174	95.3 (2020)

```

185          1,026 (2022)
186          39.9 (2020)
150          218.2 (2019)
151          96.9 (2019)
194          1,793.1 (2022)
195           3.8 (2022)
196          45.5 (2022)
197          41.2 (2022)
198           4.9 (2022)
136          2,127.2 (2020)
137          426.4 (2019)
138          122.5 (2019)
124          1,174.9 (2019)
125          2,757.4 (2019)
126           95.0 (2019)
127          243.2 (2019)
128           80.6 (2019)

```

Por suerte no hay la triple combinación ciudad-país-nombre del metro. Así tras hacer el merge no hemos de eliminar la columna Name.

```

[10]: unique_cities_metro_name = df_metro.groupby(['City', 'Country', 'Name']).size().
      ↪reset_index(name='Count')
mask = unique_cities_metro_name['Count']>1
unique_cities_metro_name[mask]

```

```

[10]: Empty DataFrame
Columns: [City, Country, Name, Count]
Index: []

```

Por otro lado, también hay varias ciudades hómonimas en mismo país:

```

[11]: unique_cities_cities = df_cities.groupby(['city_ascii', 'country']).size().
      ↪reset_index(name='Count')
mask_cities_cities = unique_cities_cities['Count']>1
unique_cities_cities[mask_cities_cities]

```

```

[11]:
   city_ascii  country  Count
55    Abasolo    Mexico     2
56     Abay    Kazakhstan     2
80   Aberdeen  United States     4
97   Abington  United States     2
169   Acatlan    Mexico     2
...         ...         ...
42236 Zhangjiazhuang    China     2
42239   Zhangping    China     2
42262  Zheleznogorsk    Russia     2
42282   Zhijiang    China     2

```

42301	Zhongshan	China	2
-------	-----------	-------	---

[1452 rows x 3 columns]

Pero de todas esas, solo nos competen las que tienen metro:

```
[12]: pd.merge(unique_cities_metro, unique_cities_cities[mask_cities_cities],
               left_on=["City", "Country"],
               right_on=["city_ascii", "country"])[['City', 'Country', 'Count_y']]
```

```
[12]:
```

	City	Country	Count_y
0	Changsha	China	2
1	Cleveland	United States	3
2	Dongguan	China	2
3	Fuzhou	China	2
4	Gwangju	South Korea	2
5	Jaipur	India	2
6	Miami	United States	2
7	Suzhou	China	2
8	Taizhou	China	2
9	Wuxi	China	2

Vamos a suponer en cada conflicto que la ciudad con metro será la de mayor población. Es decir, vuelvo a guardar como ya hice en `unique_cities_metro` los índices del dataset original. Vuelvo a mergeear. Y en cada registro (conflicto) guardo índices de ciudades mas pequeñas. Las que finalmente elimino del dataset de ciudades.

```
[13]: unique_cities_cities = df_cities.groupby(['city_ascii', 'country']).agg({'lat': λ
    ↪ lambda x: x.index.tolist()}).reset_index()
unique_cities_cities.rename(columns={'lat': 'Original indexes'}, inplace=True)
unique_cities_cities['Count'] = [len(k) for k in unique_cities_cities['Original_
    ↪ indexes']]
mask_cities_cities = unique_cities_cities['Count']>1
unique_cities_cities[mask_cities_cities]
```

```
[13]:
```

	city_ascii	country	Original indexes	Count
55	Abasolo	Mexico	[7800, 36261]	2
56	Abay	Kazakhstan	[22412, 38952]	2
80	Aberdeen	United States	[20999, 26822, 29107, 29794]	4
97	Abington	United States	[14521, 28967]	2
169	Acatlan	Mexico	[22509, 24113]	2
...
42236	Zhangjiazhuang	China	[15064, 25281]	2
42239	Zhangping	China	[2782, 44147]	2
42262	Zheleznogorsk	Russia	[5920, 12556]	2
42282	Zhijiang	China	[1574, 5287]	2
42301	Zhongshan	China	[275, 21491]	2

[1452 rows x 4 columns]

```
[14]: df_merge = pd.merge(unique_cities_metro,
    ↪ unique_cities_cities[mask_cities_cities],
        left_on=["City", "Country"],
        right_on=["city_ascii", "country"])
df_merge[['City', 'Country', 'Count_y', 'Original indexes_y']]
```

```
[14]:
```

	City	Country	Count_y	Original indexes_y
0	Changsha	China	2	[151, 1165]
1	Cleveland	United States	3	[474, 13250, 37437]
2	Dongguan	China	2	[38, 12924]
3	Fuzhou	China	2	[186, 207]
4	Gwangju	South Korea	2	[527, 2273]
5	Jaipur	India	2	[254, 38513]
6	Miami	United States	2	[100, 34184]
7	Suzhou	China	2	[125, 172]
8	Taizhou	China	2	[135, 383]
9	Wuxi	China	2	[241, 1026]

```
[15]: df_merge.shape[0] , df_merge[['Count_y']].sum().values[0] # cities unique vs
    ↪ cities total
```

```
[15]: (10, 21)
```

```
[16]: indexes = list(chain.from_iterable(df_merge['Original indexes_y']))
df_cities.iloc[indexes ,:][['country', 'city_ascii', 'population']]
```

```
[16]:
```

	country	city_ascii	population
151	China	Changsha	4766296.0
1165	China	Changsha	717700.0
474	United States	Cleveland	1683059.0
13250	United States	Cleveland	72589.0
37437	United States	Cleveland	11285.0
38	China	Dongguan	10646000.0
12924	China	Dongguan	75135.0
186	China	Fuzhou	4047200.0
207	China	Fuzhou	3671192.0
527	South Korea	Gwangju	1490092.0
2273	South Korea	Gwangju	310278.0
254	India	Jaipur	3073350.0
38513	India	Jaipur	10259.0
100	United States	Miami	5711945.0
34184	United States	Miami	12997.0
125	China	Suzhou	5352924.0
172	China	Suzhou	4330000.0
135	China	Taizhou	5031000.0

383	China	Taizhou	2162461.0
241	China	Wuxi	3245179.0
1026	China	Wuxi	853197.0

```
[17]: indexes_to_remove = list()
      for indexes in df_merge['Original indexes_y']:
          result_dict = dict(df_cities.loc[indexes, 'population'])
          max_value = max(result_dict.values())
          indexes_with_smaller_values = [key for key, value in result_dict.items() if
          ↪value < max_value]
          indexes_to_remove.append(indexes_with_smaller_values)
      indexes_to_remove = list(chain.from_iterable(indexes_to_remove))
```

```
[18]: len(indexes_to_remove) # cities to remove
```

```
[18]: 11
```

```
[19]: df_cities.shape
```

```
[19]: (44691, 6)
```

```
[20]: df_cities.drop(indexes_to_remove, axis=0, inplace=True)
```

```
[21]: df_cities.shape
```

```
[21]: (44680, 6)
```

```
[22]: # comprobación
      unique_cities_cities = df_cities.groupby(['city_ascii', 'country']).size().
      ↪reset_index(name='Count')
      mask_cities_cities = unique_cities_cities['Count'] > 1
      pd.merge(unique_cities_metro, unique_cities_cities[mask_cities_cities],
                left_on=["City", "Country"],
                right_on=["city_ascii", "country"])
```

```
[22]: Empty DataFrame
      Columns: [City, Country, Original indexes, Count_x, city_ascii, country,
      Count_y]
      Index: []
```

2.2.2 pd.merge()

```
[23]: df = pd.merge(df_metro, df_cities,
                    left_on=["City", "Country"],
                    right_on=["city_ascii", "country"],
                    how='left').drop('city_ascii', axis=1)
      df.drop('Country', axis=1, inplace=True)
```



```
[24]: df.head()
```

```
[24]:
```

	City	Name	Year opened	Last year expanded	\
0	Algiers	Algiers Metro	2011	2018	
1	Buenos Aires	Buenos Aires Underground	1913	2019	
2	Yerevan	Yerevan Metro	1981	1996	
3	Sydney	Sydney Metro	2019	-	
4	Vienna	Vienna U-Bahn	1978	2017	

	Stations	System length	Annual ridership (millions)	country	\
0	19	18.5 km (11.5 mi)	45.3 (2019)	Algeria	
1	78	56.7 km (35.2 mi)	74.0 (2020)	Argentina	
2	10	12.1 km (7.5 mi)	23.3 (2022)	Armenia	
3	13	36 km (22 mi)	16.3 (2022)	Australia	
4	98	83.3 km (51.8 mi)	459.8 (2019)	Austria	

	admin_name	population	lat	lng
0	Alger	3415811.0	36.7539	3.0589
1	Buenos Aires, Ciudad Autónoma de	16710000.0	-34.5997	-58.3819
2	Yerevan	1075800.0	40.1814	44.5144
3	New South Wales	4840600.0	-33.8678	151.2100
4	Wien	1973403.0	48.2083	16.3725

Comprobamos que las 205 ciudades con metros se corresponden 1 a 1 en el dataframe mergeado:

```
[25]: [k.shape[0] for k in [df_metro,df]]
```

```
[25]: [205, 205]
```

3 Limpieza de los datos

Estudiamos NaN

```
[26]: def df_check_nan_null(df):
    print('** NA **')
    # note: read your csv with empty cells interpreted as empty strings by
    ↪ simply setting keep_default_na=False
    na_cols = df.isnull().sum()
    print("\nNon NaN cols:\n", na_cols[na_cols != 0])

    print('\n\n** Blancos **')
    blank_cols = (df == "").sum()
    blank_cols = blank_cols[blank_cols != 0]
    print("\nNon zero cols:\n", blank_cols)

    return na_cols, blank_cols
```

```
[27]: na_cols, blank_cols = df_check_nan_null(df);
```

**** NA ****

Non NaN cols:

```
Annual ridership (millions)    16
country                        17
admin_name                    19
population                    17
lat                           17
lng                           17
dtype: int64
```

**** Blancos ****

Non zero cols:

```
Series([], dtype: int64)
```

16 registros del dataset del metro no recogen los millones de pasajeros.

3.1 Mergeo sin ASCII

Los otros campos con 17 registros muy probablemente se deban a un incorrecto mergeado:

```
[28]: def df_check_nan(df, na_cols, col_index=[1]): #[0,1]
      for k in col_index:
          col = na_cols[na_cols>0].index[k]
          print(col)
          mask = df[[col]].isnull().values
          index = df[mask].index
          display(df[mask].reset_index())
      return mask, index;
mask_nan, index_nan_merged_df = df_check_nan(df, na_cols);
```

country

	index	City	Name	Year opened \
0	10	Brasília	Federal District Metro	2001
1	16	São Paulo	São Paulo Metro	1974
2	38	Hong Kong	Mass Transit Railway	1979
3	59	Ürümqi	Ürümqi Metro	2018
4	65	Xuzhou	Xuzhou Metro	2019
5	67	Medellín	Medellín Metro	1995
6	68	Prague	Prague Metro	1974
7	94	Kanpur	Kanpur Metro	2021
8	103	Isfahan	Isfahan Urban Railway	2015
9	158	Nizhny Novgorod	Nizhny Novgorod Metro	1985
10	172	Taoyuan	Taoyuan Metro	2017
11	179	İzmir	İzmir Metro	2000
12	194	New York City	New York City Subway	1904

13	195	New York City	Staten Island Railway	1925
14	196	New York City	PATH	1908
15	199	San Francisco (Bay Area)	BART	1972
16	201	Washington, D.C.	Washington Metro	1976

	Last year expanded	Stations	System length \
0	2020	27	42.38 km (26.33 mi)
1	2021	89	104.4 km (64.9 mi)
2	2022	99	209.0 km (129.9 mi)
3	2019	21	27.615 km (17.159 mi)
4	2021	51	64.35 km (39.99 mi)
5	2012	27	31.3 km (19.4 mi)
6	2015	58	65.4 km (40.6 mi)
7	-	9	8.98 km (5.58 mi)
8	2018	20	20.2 km (12.6 mi)
9	2018	15	21.82 km (13.56 mi)
10	-	22	53.1 km (33.0 mi)
11	2014	17	20 km (12 mi)
12	2017	424	399 km (248 mi)
13	2017	21	22.5 km (14.0 mi)
14	1937	13	22.2 km (13.8 mi)
15	2020	47	191.5 km (119.0 mi)
16	2023	98	206 km (128 mi)

	Annual ridership (millions)	country	admin_name	population	lat	lng
0	42.8 (2019)	NaN	NaN	NaN	NaN	NaN
1	1,104.149 (2022)	NaN	NaN	NaN	NaN	NaN
2	1,616.30 (2021)	NaN	NaN	NaN	NaN	NaN
3	19.11 (2020)	NaN	NaN	NaN	NaN	NaN
4	20.94 (2020)	NaN	NaN	NaN	NaN	NaN
5	215.2 (2022)	NaN	NaN	NaN	NaN	NaN
6	251.4 (2020)	NaN	NaN	NaN	NaN	NaN
7	NaN	NaN	NaN	NaN	NaN	NaN
8	NaN	NaN	NaN	NaN	NaN	NaN
9	29.9 (2022)	NaN	NaN	NaN	NaN	NaN
10	28.0 (2019)	NaN	NaN	NaN	NaN	NaN
11	100 (2019)	NaN	NaN	NaN	NaN	NaN
12	1,793.1 (2022)	NaN	NaN	NaN	NaN	NaN
13	3.8 (2022)	NaN	NaN	NaN	NaN	NaN
14	45.5 (2022)	NaN	NaN	NaN	NaN	NaN
15	39.6 (2022)	NaN	NaN	NaN	NaN	NaN
16	93.0 (2022)	NaN	NaN	NaN	NaN	NaN

Hemos combinado los dataframes en función del código ASCII, pero el CSV del metro no recoge el ASCII (siempre), sino que a veces es en idioma original. Probamos a combinar estos 17 registros con el campo `city` en vez de `city_ascii` de `df_cities`. Lo cual funciona con Brasília, São Paulo, etc.

```
[29]: df_cities_orig = pd.read_csv("../dataset/worldcities.csv")
df_nan_merged = pd.merge(df[mask_nan], df_cities_orig,
                          left_on=["City"],
                          right_on=["city"],
                          how='left',
                          suffixes=('_x', ''')).drop('city', axis=1)

# Keep columns from the second DataFrame in case of homonyms
columns_to_keep = [col for col in df_nan_merged if not col.endswith('_x')]
df_nan_merged = df_nan_merged[columns_to_keep][df.columns]
```

Sobreescribimos en dataframe mergeado original los anteriores. Y comprobamos que NaN aún faltan:

```
[30]: df.iloc[index_nan_merged_df,:] = df_nan_merged
df.iloc[index_nan_merged_df,:]
```

```
[30]:
```

	City	Name	Year opened \
10	Brasília	Federal District Metro	2001
16	São Paulo	São Paulo Metro	1974
38	Hong Kong	Mass Transit Railway	1979
59	Ürümqi	Ürümqi Metro	2018
65	Xuzhou	Xuzhou Metro	2019
67	Medellín	Medellín Metro	1995
68	Prague	Prague Metro	1974
94	Kanpur	Kanpur Metro	2021
103	Isfahan	Isfahan Urban Railway	2015
158	Nizhny Novgorod	Nizhny Novgorod Metro	1985
172	Taoyuan	Taoyuan Metro	2017
179	İzmir	İzmir Metro	2000
194	New York City	New York City Subway	1904
195	New York City	Staten Island Railway	1925
196	New York City	PATH	1908
199	San Francisco (Bay Area)	BART	1972
201	Washington, D.C.	Washington Metro	1976

	Last year expanded	Stations	System length \
10	2020	27	42.38 km (26.33 mi)
16	2021	89	104.4 km (64.9 mi)
38	2022	99	209.0 km (129.9 mi)
59	2019	21	27.615 km (17.159 mi)
65	2021	51	64.35 km (39.99 mi)
67	2012	27	31.3 km (19.4 mi)
68	2015	58	65.4 km (40.6 mi)
94	-	9	8.98 km (5.58 mi)
103	2018	20	20.2 km (12.6 mi)
158	2018	15	21.82 km (13.56 mi)
172	-	22	53.1 km (33.0 mi)

179	2014	17	20 km (12 mi)
194	2017	424	399 km (248 mi)
195	2017	21	22.5 km (14.0 mi)
196	1937	13	22.2 km (13.8 mi)
199	2020	47	191.5 km (119.0 mi)
201	2023	98	206 km (128 mi)

	Annual ridership (millions)	country	admin_name	population \
10	42.8 (2019)	Brazil	Distrito Federal	3039444.0
16	1,104.149 (2022)	Brazil	São Paulo	23086000.0
38	1,616.30 (2021)	Hong Kong	NaN	7450000.0
59	19.11 (2020)	China	Xinjiang	4335017.0
65	20.94 (2020)	NaN	NaN	NaN
67	215.2 (2022)	Colombia	Antioquia	2529403.0
68	251.4 (2020)	Czechia	Praha	1335084.0
94	NaN	NaN	NaN	NaN
103	NaN	NaN	NaN	NaN
158	29.9 (2022)	NaN	NaN	NaN
172	28.0 (2019)	NaN	NaN	NaN
179	100 (2019)	Turkey	İzmir	4320519.0
194	1,793.1 (2022)	NaN	NaN	NaN
195	3.8 (2022)	NaN	NaN	NaN
196	45.5 (2022)	NaN	NaN	NaN
199	39.6 (2022)	NaN	NaN	NaN
201	93.0 (2022)	NaN	NaN	NaN

	lat	lng
10	-15.7939	-47.8828
16	-23.5500	-46.6333
38	22.3000	114.2000
59	43.8225	87.6125
65	NaN	NaN
67	6.2308	-75.5906
68	50.0875	14.4214
94	NaN	NaN
103	NaN	NaN
158	NaN	NaN
172	NaN	NaN
179	38.4200	27.1400
194	NaN	NaN
195	NaN	NaN
196	NaN	NaN
199	NaN	NaN
201	NaN	NaN

```
[31]: mask_nan, index_nan_merged_df = df_check_nan(df,na_cols);
```

country

	index	City	Name	Year opened	\
0	65	Xuzhou	Xuzhou Metro	2019	
1	94	Kanpur	Kanpur Metro	2021	
2	103	Isfahan	Isfahan Urban Railway	2015	
3	158	Nizhny Novgorod	Nizhny Novgorod Metro	1985	
4	172	Taoyuan	Taoyuan Metro	2017	
5	194	New York City	New York City Subway	1904	
6	195	New York City	Staten Island Railway	1925	
7	196	New York City	PATH	1908	
8	199	San Francisco (Bay Area)	BART	1972	
9	201	Washington, D.C.	Washington Metro	1976	

	Last year expanded	Stations	System length	\
0	2021	51	64.35 km (39.99 mi)	
1	-	9	8.98 km (5.58 mi)	
2	2018	20	20.2 km (12.6 mi)	
3	2018	15	21.82 km (13.56 mi)	
4	-	22	53.1 km (33.0 mi)	
5	2017	424	399 km (248 mi)	
6	2017	21	22.5 km (14.0 mi)	
7	1937	13	22.2 km (13.8 mi)	
8	2020	47	191.5 km (119.0 mi)	
9	2023	98	206 km (128 mi)	

	Annual ridership (millions)	country	admin_name	population	lat	lng
0	20.94 (2020)	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN
3	29.9 (2022)	NaN	NaN	NaN	NaN	NaN
4	28.0 (2019)	NaN	NaN	NaN	NaN	NaN
5	1,793.1 (2022)	NaN	NaN	NaN	NaN	NaN
6	3.8 (2022)	NaN	NaN	NaN	NaN	NaN
7	45.5 (2022)	NaN	NaN	NaN	NaN	NaN
8	39.6 (2022)	NaN	NaN	NaN	NaN	NaN
9	93.0 (2022)	NaN	NaN	NaN	NaN	NaN

3.2 Registros individuales

Ya hemos comprobado que columna ASCII no coincide con City del dataset del metro en estos 10 registros. Intentamos solventarlos individualmente.

3.2.1 Xuzhou

La web https://en.wikipedia.org/wiki/Xuzhou_Metro nos indica: - la provincia: Jiangsu Province - about 67 kilometres (42 mi) - In its first year of operation the ridership was 7.5 million (2019) - Number of stations: 54

Datos muy parecidos a los del dataset del metro:

```
[32]: df_metro[df_metro["City"] == "Xuzhou"]
```

```
[32]:      City Country      Name  Year opened Last year expanded  Stations \
65  Xuzhou   China  Xuzhou Metro      2019              2021         51
```

```
      System length Annual ridership (millions)
65  64.35 km (39.99 mi)              20.94 (2020)
```

El mandarin se transcribe muy asiduamente mal, vamos a filtrar ciudades con misma terminación:

```
[33]: mask = df_cities_orig['city_ascii'].str.endswith('uzhou')
df_cities_orig[mask]
```

```
[33]:      city city_ascii      lat      lng country iso2 iso3 admin_name \
125  Suzhou     Suzhou  33.6333  116.9683   China  CN  CHN     Anhui
172  Suzhou     Suzhou  31.3000  120.6194   China  CN  CHN    Jiangsu
178  Luzhou     Luzhou  28.8918  105.4409   China  CN  CHN    Sichuan
181  Liuzhou    Liuzhou  24.3264  109.4281   China  CN  CHN    Guangxi
186  Fuzhou     Fuzhou  27.9814  116.3577   China  CN  CHN    Jiangxi
188  Zhuzhou    Zhuzhou  27.8407  113.1469   China  CN  CHN     Hunan
190  Chuzhou    Chuzhou  32.3062  118.3115   China  CN  CHN     Anhui
207  Fuzhou     Fuzhou  26.0769  119.2917   China  CN  CHN     Fujian
257  Wuzhou     Wuzhou  23.4833  111.3167   China  CN  CHN    Guangxi
506  Huzhou     Huzhou  30.8925  120.0875   China  CN  CHN    Zhejiang
983  Quzhou     Quzhou  28.9545  118.8763   China  CN  CHN    Zhejiang
13398  Guzhou    Guzhou  25.9452  108.5238   China  CN  CHN    Guizhou
15554  Jiuzhou   Jiuzhou  39.5054  116.5642   China  CN  CHN     Hebei
30087  Luzhou     Luzhou  23.3687  114.5196   China  CN  CHN   Guangdong
```

```
      capital population      id
125      NaN   5352924.0  1156871297
172    minor   4330000.0  1156029196
178    minor   4218427.0  1156582079
181    minor   4157934.0  1156360785
186      NaN   4047200.0  1156915325
188    minor   4020800.0  1156041962
190      NaN   3987054.0  1156036420
207    admin   3671192.0  1156188037
257    minor   3061100.0  1156620133
506    minor   1558826.0  1156335543
983      NaN    902767.0  1156180376
13398    NaN     70098.0  1156435005
15554    NaN     49616.0  1156799658
30087    NaN     15890.0  1156708150
```

Vemos que el segundo resultado también es el único de la misma provincia que lo de Wikipedia:

```
[34]: mask2 = df_cities_orig["admin_name"] == "Jiangsu"
      match = df_cities_orig[mask*mask2]
      match
```

```
[34]:      city city_ascii  lat      lng country iso2 iso3 admin_name capital \
172  Suzhou      Suzhou 31.3  120.6194  China  CN  CHN   Jiangsu  minor

      population      id
172  4330000.0  1156029196
```

Ya teníamos datos de una ciudad homónima del mismo país, pero comprobamos que ni Name (del sistema de metro de la ciudad) ni admin_name, así como otros campos, no son de la ciudad que nos compete.

```
[35]: df.loc[df["City"] == "Suzhou", :]
```

```
[35]:      City      Name  Year opened Last year expanded  Stations \
55  Suzhou  Suzhou Rail Transit      2012              2021      154

      System length Annual ridership (millions) country admin_name \
55  208.2 km (129.4 mi)      308.57 (2020)  China      Anhui

      population      lat      lng
55  5352924.0  33.6333  116.9683
```

Como el nombre de la ciudad y su país son iguales en ambos registros, y no hay tercer campo con el que filtrar la unión de datasets, pues es en vano simplemente corregir el nombre de la ciudad para después mergear. Directamente añadido valores NaN:

```
[36]: cols_cities = "country,admin_name,population,lat,lng".split(",")
```

```
[37]: df.loc[df["City"] == "Xuzhou", cols_cities] = match[cols_cities].values
```

Comparamos ambas ciudades: “Suzhou, Anhui, China” vs “Xuzhou, Jiangsu, China”. Como latitud y longitud denotan, están muy próximas: según Google Maps a menos de 100 km ([link](#)).

```
[38]: df[df["City"].isin(["Xuzhou", "Suzhou"])]
```

```
[38]:      City      Name  Year opened Last year expanded  Stations \
55  Suzhou  Suzhou Rail Transit      2012              2021      154
65  Xuzhou      Xuzhou Metro      2019              2021       51

      System length Annual ridership (millions) country admin_name \
55  208.2 km (129.4 mi)      308.57 (2020)  China      Anhui
65  64.35 km (39.99 mi)      20.94 (2020)  China      Jiangsu

      population      lat      lng
55  5352924.0  33.6333  116.9683
65  4330000.0  31.3000  120.6194
```


3.2.2 Kanpur

Los datos del metro de Wikipedia coinciden aprox. con los de nuestro dataset:
https://en.wikipedia.org/wiki/Kanpur_Metro

```
[39]: df_metro[df_metro["City"] == "Kanpur"]
```

```
[39]:      City Country      Name  Year opened Last year expanded  Stations \
94  Kanpur   India  Kanpur Metro      2021                -         9

      System length Annual ridership (millions)
94  8.98 km (5.58 mi)                      NaN
```

Parece que lo han transcrito mal. Añadiendo una “h” tras la K. Entre Khanpur y Khānpur cogemos la ciudad claramente mas grande.

```
[40]: mask = df_cities_orig['city_ascii'].str.endswith('npur') &
      df_cities_orig['city_ascii'].str.startswith('K')
      df_cities_orig[mask]
```

```
[40]:      city      city_ascii      lat      lng  country iso2 iso3 \
3946   Khanpur      Khanpur  28.6453  70.6567  Pakistan  PK  PAK
10027  Khānpur      Khanpur  25.8572  85.9330    India   IN  IND
10613  Kanchanpur  Kanchanpur  25.6636  85.2703    India   IN  IND
31014  Kaleyānpur  Kaleyānpur  26.4297  84.9327    India   IN  IND
34532  Kishunpur   Kishunpur  25.3272  87.7173    India   IN  IND
35143  Khānjahānpur  Khanjahanpur  25.6055  86.0927    India   IN  IND
36694  Kamānpur     Kamanpur  18.6667  79.5000    India   IN  IND
37944  Kishunpur   Kishunpur  25.7947  86.8237    India   IN  IND
38149  Kalyānpur   Kalyanpur  26.4802  84.1789    India   IN  IND
40066  Kanchanpur  Kanchanpur  24.6096  84.2361    India   IN  IND
```

```
      admin_name capital  population      id
3946      Punjab     NaN    160308.0  1586169401
10027     Bihār     NaN     12066.0  1356667678
10613     Bihār     NaN      8616.0  1356097827
31014     Bihār     NaN    13704.0  1356786428
34532     Bihār     NaN    11237.0  1356155445
35143     Bihār     NaN    10899.0  1356018473
36694  Andhra Pradesh     NaN    11048.0  1356269111
37944     Bihār     NaN      9963.0  1356145353
38149     Bihār     NaN     9802.0  1356695468
40066     Bihār     NaN     9758.0  1356197058
```

```
[41]: mask = df_cities_orig['city'] == 'Khanpur'
      match = df_cities_orig[mask]
      match
```

```
[41]:      city city_ascii      lat      lng  country iso2 iso3 admin_name \
3946  Khanpur      Khanpur  28.6453  70.6567  Pakistan  PK  PAK      Punjab

      capital  population      id
3946      NaN    160308.0  1586169401
```

Pero hay error en el país:

```
[42]: match = match.copy()
match.loc[:, 'country'] = "India"
match
```

```
[42]:      city city_ascii      lat      lng country iso2 iso3 admin_name \
3946  Khanpur      Khanpur  28.6453  70.6567   India  PK  PAK      Punjab

      capital  population      id
3946      NaN    160308.0  1586169401
```

```
[43]: df.loc[df["City"] == "Khanpur", cols_cities] = match[cols_cities].values
```

3.2.3 Isfahan

La web del metro (https://en.wikipedia.org/wiki/Isfahan_Metro) nos hipervincula con la Wikipedia de la ciudad, ésta indica que la ciudad tiene 2.2 millones de habitantes

```
[44]: df_metro[df_metro["City"] == "Isfahan"]
```

```
[44]:      City Country      Name  Year opened Last year expanded \
103  Isfahan      Iran  Isfahan Urban Railway      2015      2018

      Stations      System length Annual ridership (millions)
103         20  20.2 km (12.6 mi)                        NaN
```

Por tanto, la ciudad es Esfahan:

```
[45]: mask = df_cities_orig['city_ascii'].str.endswith('ahan')
df_cities_orig[mask]
```

```
[45]:      city      city_ascii      lat      lng      country \
377      Eşfahān      Esfahan  32.6447  51.6675      Iran
4979      Behbahān      Behbahan  30.5958  50.2417      Iran
6862      Pinagkaisahan      Pinagkaisahan  14.5229  121.0555      Philippines
7478      Pinyahan      Pinyahan  14.6400  121.0461      Philippines
7608      Ardahan      Ardahan  41.1111  42.7022      Turkey
8176      Mukdahan      Mukdahan  16.5431  104.7228      Thailand
13012      Pinamungahan      Pinamungahan  10.2667  123.5833      Philippines
18238      Simbahan      Simbahan  6.3000  120.5833      Philippines
19399      Diplahan      Diplahan  7.6911  122.9853      Philippines
20752      Nāhan      Nahan  30.5500  77.3000      India
```

25831		Hanahan	Hanahan	32.9302	-80.0027	United States
31755		Dargahān	Dargahan	26.9636	56.0622	Iran
38115	Phibun Mangsahan	Phibun Mangsahan		15.2482	105.2296	Thailand
38369	Kūcheṣfahān	Kuchesfahan		37.2783	49.7728	Iran
41817	Dānesfahān	Danesfahan		35.8108	49.7422	Iran
42529	Harahan	Harahan		29.9374	-90.2031	United States

	iso2	iso3	admin_name	capital	population	id
377	IR	IRN	Eṣfahān	admin	2219343.0	1364023865
4979	IR	IRN	Khūzestān	minor	122604.0	1364393434
6862	PH	PHL	Makati	NaN	57343.0	1608216406
7478	PH	PHL	Quezon	NaN	28129.0	1608055244
7608	TR	TUR	Ardahan	admin	42226.0	1792379425
8176	TH	THA	Mukdahan	admin	33102.0	1764994534
13012	PH	PHL	Cebu	NaN	75131.0	1608414270
18238	PH	PHL	Sulu	NaN	36374.0	1608320501
19399	PH	PHL	Zamboanga Sibugay	NaN	32585.0	1608903309
20752	IN	IND	Himāchal Pradesh	NaN	28899.0	1356417528
25831	US	USA	South Carolina	NaN	20381.0	1840014256
31755	IR	IRN	Hormozgān	NaN	14525.0	1364735795
38115	TH	THA	Ubon Ratchathani	minor	10842.0	1764591980
38369	IR	IRN	Gīlān	NaN	10026.0	1364862389
41817	IR	IRN	Qazvīn	NaN	9434.0	1364115128
42529	US	USA	Louisiana	NaN	9137.0	1840013997

```
[46]: mask = df_cities_orig['city_ascii'] == 'Esfahan'
      match = df_cities_orig[mask]
      match
```

```
[46]:      city city_ascii      lat      lng country iso2 iso3 admin_name \
377  Eṣfahān    Esfahan  32.6447  51.6675    Iran   IR   IRN    Eṣfahān

      capital population      id
377    admin    2219343.0  1364023865
```

```
[47]: df.loc[df["City"] == "Isfahan", cols_cities] = match[cols_cities].values
```

3.2.4 Nizhny Novgorod

```
[48]: df_metro[df_metro["City"] == "Nizhny Novgorod"]
```

```
[48]:      City Country      Name  Year opened \
158  Nizhny Novgorod  Russia  Nizhny Novgorod Metro    1985

      Last year expanded  Stations      System length \
158                2018         15  21.82 km (13.56 mi)
```

Annual ridership (millions)
158 29.9 (2022)

```
[49]: mask = df_cities_orig['city_ascii'].str.endswith('Novgorod')
df_cities_orig[mask]
```

```
[49]:
```

	city	city_ascii	lat	lng	country	iso2	iso3	\
606	Nizhniy Novgorod	Nizhniy Novgorod	56.3269	44.0075	Russia	RU	RUS	
2959	Velikiy Novgorod	Velikiy Novgorod	58.5210	31.2758	Russia	RU	RUS	

	admin_name	capital	population	id
606	Nizhegorodskaya Oblast'	admin	1264075.0	1643012126
2959	Novgorodskaya Oblast'	admin	222594.0	1643774241

```
[50]: mask = df_cities_orig['city_ascii'] == 'Nizhniy Novgorod'
match = df_cities_orig[mask]
match
```

```
[50]:
```

	city	city_ascii	lat	lng	country	iso2	iso3	\
606	Nizhniy Novgorod	Nizhniy Novgorod	56.3269	44.0075	Russia	RU	RUS	

	admin_name	capital	population	id
606	Nizhegorodskaya Oblast'	admin	1264075.0	1643012126

```
[51]: df.loc[df["City"] == "Nizhny Novgorod", cols_cities] = match[cols_cities].values
```

3.2.5 Taoyuan

```
[52]: df_metro[df_metro["City"] == "Taoyuan"]
```

```
[52]:
```

	City	Country	Name	Year opened	Last year expanded	Stations	\
172	Taoyuan	Taiwan	Taoyuan Metro	2017	-	22	

	System length	Annual ridership (millions)
172	53.1 km (33.0 mi)	28.0 (2019)

```
[53]: mask = df_cities_orig['iso2'] == "TW"
df_cities_orig[mask]
```

```
[53]:
```

	city	city_ascii	lat	lng	country	iso2	iso3	admin_name	\
47	Taipei	Taipei	25.0375	121.5625	Taiwan	TW	TWN	Taipei	
262	Taichung	Taichung	24.1439	120.6794	Taiwan	TW	TWN	Taichung	
292	Kaohsiung	Kaohsiung	22.6150	120.2975	Taiwan	TW	TWN	Kaohsiung	
437	Tainan	Tainan	22.9833	120.1833	Taiwan	TW	TWN	Tainan	
486	Zhongli	Zhongli	24.9650	121.2168	Taiwan	TW	TWN	Taoyuan	
...	
41269	Fuli	Fuli	23.1333	121.2833	Taiwan	TW	TWN	Hualien	
41356	Xinpi	Xinpi	22.4880	120.5814	Taiwan	TW	TWN	Pingtung	

42650	Nanzhuang	Nanzhuang	24.5699	121.0157	Taiwan	TW	TWN	Miaoli
43176	Beibu	Beibu	24.6639	121.0681	Taiwan	TW	TWN	Hsinchu
44143	Jianshi	Jianshi	24.5761	121.3081	Taiwan	TW	TWN	Hsinchu

	capital	population	id
47	primary	9078000.0	1158881289
262	admin	3033885.0	1158689622
292	admin	2733566.0	1158331334
437	admin	1874686.0	1158061376
486	NaN	1632616.0	1158025380
...
41269	NaN	9681.0	1158634303
41356	NaN	9540.0	1158537415
42650	NaN	9029.0	1158569080
43176	NaN	8647.0	1158656385
44143	NaN	9532.0	1158994660

[159 rows x 11 columns]

```
[54]: df_cities_orig.loc[mask & df_cities_orig['city_ascii'].str.startswith('Tao')]
```

```
[54]:
```

	city	city_ascii	lat	lng	country	iso2	iso3	\
1701	Taoyuan District	Taoyuan District	24.9913	121.3143	Taiwan	TW	TWN	

	admin_name	capital	population	id
1701	Taoyuan	admin	443273.0	1158127875

```
[55]: mask = df_cities_orig['city_ascii'] == 'Taoyuan District'
match = df_cities_orig[mask]
match
```

```
[55]:
```

	city	city_ascii	lat	lng	country	iso2	iso3	\
1701	Taoyuan District	Taoyuan District	24.9913	121.3143	Taiwan	TW	TWN	

	admin_name	capital	population	id
1701	Taoyuan	admin	443273.0	1158127875

```
[56]: df.loc[df["City"] == "Taoyuan", cols_cities] = match[cols_cities].values
```

3.2.6 New York City

```
[57]: mask = df_cities_orig['city_ascii'].str.startswith('New York')
match = df_cities_orig[mask]
match
```

```
[57]:
```

	city	city_ascii	lat	lng	country	iso2	iso3	admin_name	\
11	New York	New York	40.6943	-73.9249	United States	US	USA	New York	

	capital	population	id
11	NaN	18972871.0	1840034016

```
[58]: df.loc[df["City"] == "New York City", cols_cities] = match[cols_cities].values
```

3.2.7 San Francisco (Bay Area)

```
[59]: mask = df_cities_orig['iso2'] == "US"
match = df_cities_orig.loc[mask & df_cities_orig['city_ascii'].str.
        ↪startswith('San Francisco')]
match
```

```
[59]:
```

	city	city_ascii	lat	lng	country	iso2	iso3	\
237	San Francisco	San Francisco	37.7558	-122.4449	United States	US	USA	

	admin_name	capital	population	id
237	California	NaN	3290197.0	1840021543

```
[60]: df.loc[df["City"] == "San Francisco (Bay Area)", cols_cities] =
        ↪match[cols_cities].values
```

3.2.8 Washington, D.C.

```
[61]: mask = df_cities_orig['admin_name'] == "District of Columbia"
match = df_cities_orig.loc[mask & df_cities_orig['city_ascii'].str.
        ↪startswith('Washington')]
match
```

```
[61]:
```

	city	city_ascii	lat	lng	country	iso2	iso3	\
149	Washington	Washington	38.9047	-77.0163	United States	US	USA	

	admin_name	capital	population	id
149	District of Columbia	primary	4810669.0	1840006060

```
[62]: df.loc[df["City"] == "Washington, D.C.", cols_cities] = match[cols_cities].
        ↪values
```

3.3 Comprobación

Nos faltan 3 identificadores municipales. Pero su relevancia es escasa.

```
[63]: na_cols, blank_cols = df_check_nan_null(df);
```

```
** NA **
```

Non NaN cols:

Annual ridership (millions)	16
-----------------------------	----

```
admin_name          3
dtype: int64
```

```
** Blancos **
```

```
Non zero cols:
Series([], dtype: int64)
```

```
[64]: df_check_nan(df, na_cols);
```

```
admin_name

   index  City      Name  Year opened  Last year expanded \
0     38  Hong Kong  Mass Transit Railway      1979      2022
1    129   Almaty    Almaty Metro      2011      2022
2    163  Singapore  Mass Rapid Transit      1987      2022

   Stations  System length  Annual ridership (millions)  country \
0         99  209.0 km (129.9 mi)      1,616.30 (2021)  Hong Kong
1         11   13.4 km (8.3 mi)      17.1 (2022)  Kazakhstan
2        134  230.2 km (143.0 mi)      766.5 (2021)  Singapore

   admin_name  population    lat    lng
0         NaN   7450000.0  22.3000  114.2000
1         NaN   1916822.0  43.2775   76.8958
2         NaN   5983000.0   1.3000  103.8000
```

3.4 Corrección de campos

```
[65]: def check_no_numeric_rows(df=df, col='Ridership (millions)'):
        non_float_rows = df[~pd.to_numeric(df[col], errors="coerce").notnull()]
        return non_float_rows
```

3.4.1 Solo NaN como tal

El campo Last year expanded registra - si no ha habido expansión de la red de metro alguna desde su construcción.

```
[66]: col = "Last year expanded"
        check_no_numeric_rows(col=col).head(3)
```

```
[66]:      City      Name  Year opened  Last year expanded  Stations \
3    Sydney  Sydney Metro      2019      -      13
6    Dhaka   Dhaka Metro Rail      2022      -      9
29  Dongguan  Dongguan Rail Transit      2016      -      15

      System length  Annual ridership (millions)  country \
3      36 km (22 mi)      16.3 (2022)  Australia
```

6	11.7 km (7.3 mi)	NaN	Bangladesh
29	37.7 km (23.4 mi)	35.06 (2020)	China

	admin_name	population	lat	lng
3	New South Wales	4840600.0	-33.8678	151.2100
6	Dhaka	18627000.0	23.7639	90.3889
29	Guangdong	10646000.0	23.0475	113.7493

```
[67]: df[col].replace("-", np.nan, inplace=True)
df[col] = df[col].astype(float)
check_no_numeric_rows(col=col).head(3)
```

```
[67]:
```

	City	Name	Year opened	Last year expanded	\
3	Sydney	Sydney Metro	2019	NaN	
6	Dhaka	Dhaka Metro Rail	2022	NaN	
29	Dongguan	Dongguan Rail Transit	2016	NaN	

	Stations	System length	Annual ridership (millions)	country	\
3	13	36 km (22 mi)	16.3 (2022)	Australia	
6	9	11.7 km (7.3 mi)	NaN	Bangladesh	
29	15	37.7 km (23.4 mi)	35.06 (2020)	China	

	admin_name	population	lat	lng
3	New South Wales	4840600.0	-33.8678	151.2100
6	Dhaka	18627000.0	23.7639	90.3889
29	Guangdong	10646000.0	23.0475	113.7493

Hay 22 metros sin ampliaciones

```
[68]: df_check_nan_null(df);
```

```
** NA **
```

Non NaN cols:

```
Last year expanded      22
Annual ridership (millions) 16
admin_name              3
dtype: int64
```

```
** Blancos **
```

Non zero cols:

```
Series([], dtype: int64)
```

Campo de año de construcción siempre aparece:

```
[69]: check_no_numeric_rows(col="Year opened").head(3)
```



```
[69]: Empty DataFrame
Columns: [City, Name, Year opened, Last year expanded, Stations, System length,
Annual ridership (millions), country, admin_name, population, lat, lng]
Index: []
```

3.4.2 Uso del metro: ridership

```
[70]: col = 'Annual ridership (millions)'
df[['Ridership (millions)', 'Ridership Year']] = df[col].str.extract(r'^(\d+\.\d+)\s+\((\d+)\)$')
df.drop(col, axis=1, inplace=True)

col = 'Ridership (millions)'
df[col] = df[col].astype(float)

df.head(3)
```

```
[70]:
```

	City	Name	Year opened	Last year expanded	\
0	Algiers	Algiers Metro	2011	2018.0	
1	Buenos Aires	Buenos Aires Underground	1913	2019.0	
2	Yerevan	Yerevan Metro	1981	1996.0	

	Stations	System length	country	admin_name	\
0	19	18.5 km (11.5 mi)	Algeria	Alger	
1	78	56.7 km (35.2 mi)	Argentina	Buenos Aires, Ciudad Autónoma de	
2	10	12.1 km (7.5 mi)	Armenia	Yerevan	

	population	lat	lng	Ridership (millions)	Ridership Year
0	3415811.0	36.7539	3.0589	45.3	2019
1	16710000.0	-34.5997	-58.3819	74.0	2020
2	1075800.0	40.1814	44.5144	23.3	2022

```
[71]: col = 'Ridership (millions)'
df[col] = df[col].astype(float)
```

```
[72]: check_no_numeric_rows().head()
```

```
[72]:
```

	City	Name	Year opened	Last year expanded	Stations	\
6	Dhaka	Dhaka Metro Rail	2022	NaN	9	
15	Salvador	Salvador Metro	2014	2018.0	19	
16	São Paulo	São Paulo Metro	1974	2021.0	89	
22	Beijing	Beijing Subway	1971	2023.0	370	
26	Chengdu	Chengdu Metro	2010	2020.0	284	

	System length	country	admin_name	population	lat	lng	\
6	11.7 km (7.3 mi)	Bangladesh	Dhaka	18627000.0	23.7639	90.3889	
15	32.5 km (20.2 mi)	Brazil	Bahia	2886698.0	-12.9747	-38.4767	

16	104.4 km (64.9 mi)	Brazil	São Paulo	23086000.0	-23.5500	-46.6333
22	785.7 km (488.2 mi)	China	Beijing	18522000.0	39.9040	116.4075
26	518.5 km (322.2 mi)	China	Sichuan	14645000.0	30.6600	104.0633

	Ridership (millions)	Ridership Year
6	NaN	NaN
15	NaN	NaN
16	NaN	NaN
22	NaN	NaN
26	NaN	NaN

Añado el uso por habitante de cada ciudad:

```
[73]: df['Ridership per capita'] = df['Ridership (millions)']*1e6 / df['population']
```

3.4.3 Longitud de vías

```
[74]: col = 'System length'
df[col] = df[col].str.extract(r'^([\d.]+)')
df[col] = df[col].astype(float)
df.tail(3)
```

```
[74]:
```

	City	Name	Year opened	Last year expanded	Stations	\
202	Tashkent	Tashkent Metro	1977	2023.0	48	
203	Caracas	Caracas Metro	1983	2015.0	49	
204	Hanoi	Hanoi Metro	2021	NaN	12	

	System length	country	admin_name	population	lat	\
202	59.1	Uzbekistan	Toshkent	2571668.0	41.3111	
203	67.2	Venezuela	Distrito Capital	2245744.0	10.4806	
204	13.1	Vietnam	Hà Nội	8246600.0	21.0283	

	lng	Ridership (millions)	Ridership Year	Ridership per capita
202	69.2797	136.7	2022	53.156162
203	-66.9036	NaN	NaN	NaN
204	105.8542	NaN	NaN	NaN

```
[75]: check_no_numeric_rows(col=col)
```

```
[75]: Empty DataFrame
Columns: [City, Name, Year opened, Last year expanded, Stations, System length,
country, admin_name, population, lat, lng, Ridership (millions), Ridership Year,
Ridership per capita]
Index: []
```

4 Análisis de los datos

Comprobamos que los campos enteros ó decimales son así:

```
[76]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 205 entries, 0 to 204
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   City                  205 non-null   object
1   Name                  205 non-null   object
2   Year opened           205 non-null   int64
3   Last year expanded    183 non-null   float64
4   Stations              205 non-null   int64
5   System length         205 non-null   float64
6   country               205 non-null   object
7   admin_name            202 non-null   object
8   population            205 non-null   float64
9   lat                   205 non-null   float64
10  lng                   205 non-null   float64
11  Ridership (millions)  140 non-null   float64
12  Ridership Year        140 non-null   object
13  Ridership per capita  140 non-null   float64
dtypes: float64(7), int64(2), object(5)
memory usage: 24.0+ KB
```

4.1 Distribución de campos numéricos

Estudiamos la estadística descriptiva de las variables continuas. Si título en rojo entonces menos del 95% de valores se encuentran entre percentiles 2.5 y 97.5. Entre corchetes se muestra: [percentil 5, media, percentil 95].

```
[77]: def box_violin_plot(df, cols):
    fig_cols = 5
    fig_rows = int(np.ceil(len(cols)/fig_cols))

    fig = plt.figure(figsize=(13, 4*fig_rows))
    fig.subplots_adjust(hspace=0.4, wspace=0.4)

    for i,col in enumerate(cols):
        x = df[col].dropna().values

        percentiles = [np.percentile(x, k) for k in [5,50,95]]
        percent_range_5_95 = (1 - ((sum(x<percentiles[0]) +
↪sum(x>percentiles[-1]))/len(x)))*100
        color = 'k'
        if percent_range_5_95< 95: color='r'

        ax = fig.add_subplot(fig_rows, fig_cols, i+1)
```

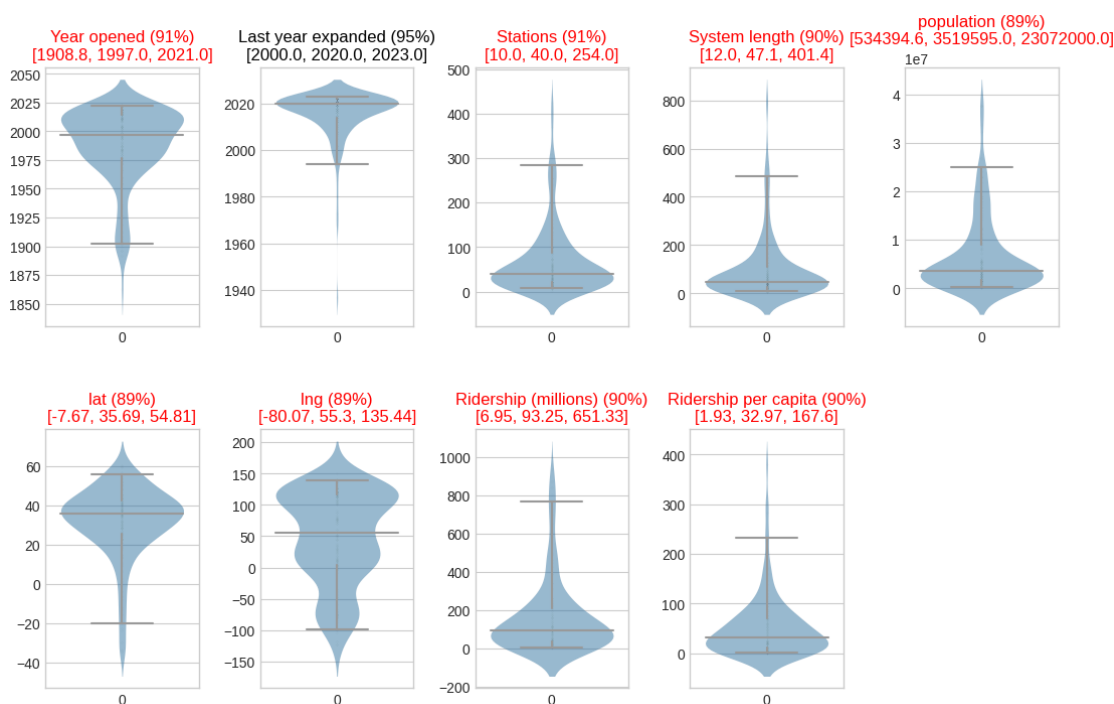
```

sns.boxplot(x, showfliers=False, showbox=False, whis=[2.5,97.5],
color='w')
sns.violinplot(x, inner='point', linewidth=.01)
plt.setp(ax.collections, alpha=.5)
ax.set_title(f"{col} ({percent_range_5_95:.0f}%) \n{round(g,2) for
g in percentiles ]}", color=color)
return fig

```

```
[78]: cols_num = df.select_dtypes(include=[np.number]).columns
```

```
[79]: with warnings.catch_warnings():
    box_violin_plot(df, cols_num);
```



Se observa que: - el boom de los metros se inicio en la década de los 50 - la gran mayoría de redes metropolitanas no llegan al centenar de estaciones - una silueta distributiva casi idéntica al n° de estaciones es la de la longitud (ver correlaciones en siguiente apartado) - la mayoría de ciudades no llegan a la decena de millones de habitantes - el uso per capita toma rangos tan variado como case 2 para el percentil 5 y menos de 165 para el percentil 95, es decir, unas $165/5=33$ veces mas. Muy dispar.

Estos gráficos nos ayudan también a identificar valores extremos (outliers). Por ejemplo que el n° de estaciones o la población tuviera valores negativos, latitudes fuera de rango $[-90,90]$, etc. Nada de esto sucede si nos fijamos en los boxplots, aunque podemos codificarlo para cerciorarnos:

```
[80]: df[cols_num].describe()
```

```
[80]:
```

	Year opened	Last year expanded	Stations	System length	\
count	205.000000	183.000000	205.000000	205.000000	
mean	1988.907317	2015.918033	69.243902	99.689698	
std	32.145902	10.565799	76.443360	132.231555	
min	1863.000000	1937.000000	2.000000	4.100000	
25%	1977.000000	2014.000000	21.000000	28.000000	
50%	1997.000000	2020.000000	40.000000	47.100000	
75%	2014.000000	2022.000000	88.000000	109.000000	
max	2023.000000	2023.000000	424.000000	795.500000	

	population	lat	lng	Ridership (millions)	\
count	2.050000e+02	205.000000	205.000000	140.000000	
mean	6.776130e+06	31.989243	48.288138	161.934500	
std	7.816394e+06	18.127545	73.158234	196.443906	
min	3.538400e+04	-34.599700	-123.100000	0.400000	
25%	1.683059e+06	25.286700	4.500000	36.500000	
50%	3.519595e+06	35.689700	55.297200	93.250000	
75%	8.911000e+06	42.318800	116.407500	209.207500	
max	3.773200e+07	60.170800	151.210000	935.200000	

	Ridership per capita
count	140.000000
mean	52.219741
std	61.779523
min	0.049944
25%	11.294210
50%	32.967561
75%	69.845015
max	380.850837

4.2 Correlaciones

```
[81]: def nor(x):
    range_vals = max(x) - min(x)
    if not range_vals: return x
    return (x - min(x))/range_vals

def plot_corr(df, cols, annotbool=True, figsize=(5, 5)):
    corr = df[cols].corr()*100

    plt.figure(figsize=figsize)
    sns.heatmap(corr, annot=annotbool, mask=np.triu(corr), fmt=".1f",
    cmap='RdBu', vmin=-100, vmax=100)
    return corr;
```


```
[82]: df_norm = copy.deepcopy(df.loc[:,cols_num])
```

```

for col in cols_num:
    df_norm[col] = nor(df_norm[col])

corr = plot_corr(df_norm, cols_num);

```



Como los violinplots ya dejaban entrever, el número de estaciones es fuertemente correlativo con la longitud del metro (93.1%).

El número de viajeros (Ridership (millions)) se correlaciona con los anteriores dos campos al 73 y 67% respectivamnete.

Estas tres correlaciones son lógicas: - a mas estaciones mas distancia total - a mas estaciones (y mayores distancias abarcables) pues mas viajeros se benefician del sistema ferroviario subterráneo

```

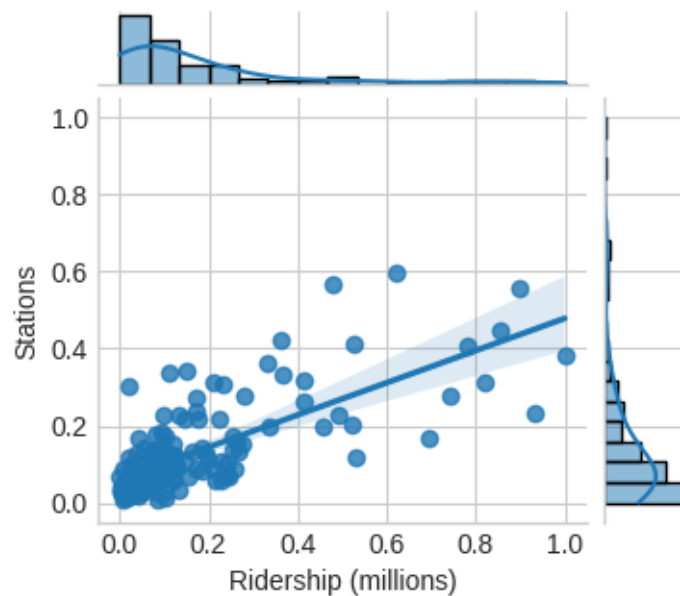

[83]: for k in ['System length', 'Ridership (millions)']:
        g = sns.jointplot(data=df_norm, x=k, y='Stations', kind='reg')


```

```

g.fig.set_size_inches((3.5, 3))
g = sns.jointplot(data=df_norm, x='System length', y='Ridership (millions)',
kind='reg')
g.fig.set_size_inches((3.5, 3));

```





El resto de variables numérica están poco correlacionadas. A resaltar el 42% entre la longitud geográfica `lng` y el año de inauguración, que como se grafica en siguiente apartado, denota la incipiente construcción de líneas de metro nuevas en China y el lejano oriente en general. A como ciudades grandes requieren mas estaciones y cobertura espacial (y viceversa en pequeñas), es decir, correlaciones superiores al 40% entre `population` y `estations` como entre `population` y `System length`.

Si eliminamos los outliers (aquellos que se alejan mas de 3 desviaciones estándar de la media) las regresiones anteriores (excepto una) no se darán tan fuertes, como se denota al rechequear las correlaciones:

```
[84]: def remove_outliers(col, limit=3):
      mn = np.nanmean(col)
      out = limit * np.nanstd(col)
      mask = (col < (mn - out)) | (col > (mn + out))
      return col[~mask]
```

```
[85]: df_norm_no_outliers = df_norm.copy()
      df_norm_no_outliers = df_norm_no_outliers.apply(remove_outliers)
```

Los outliers no se han eliminado, sino reemplazados por NaN:

```
[86]: df_norm.shape[0], df_norm_no_outliers.shape[0]
```

```
[86]: (205, 205)
```

```
[87]: df_check_nan_null(df_norm_no_outliers);
```



```
** NA **
```

```
Non NaN cols:
```

Year opened	2
Last year expanded	26
Stations	5
System length	6
population	5
lat	6
Ridership (millions)	70
Ridership per capita	68

```
dtype: int64
```

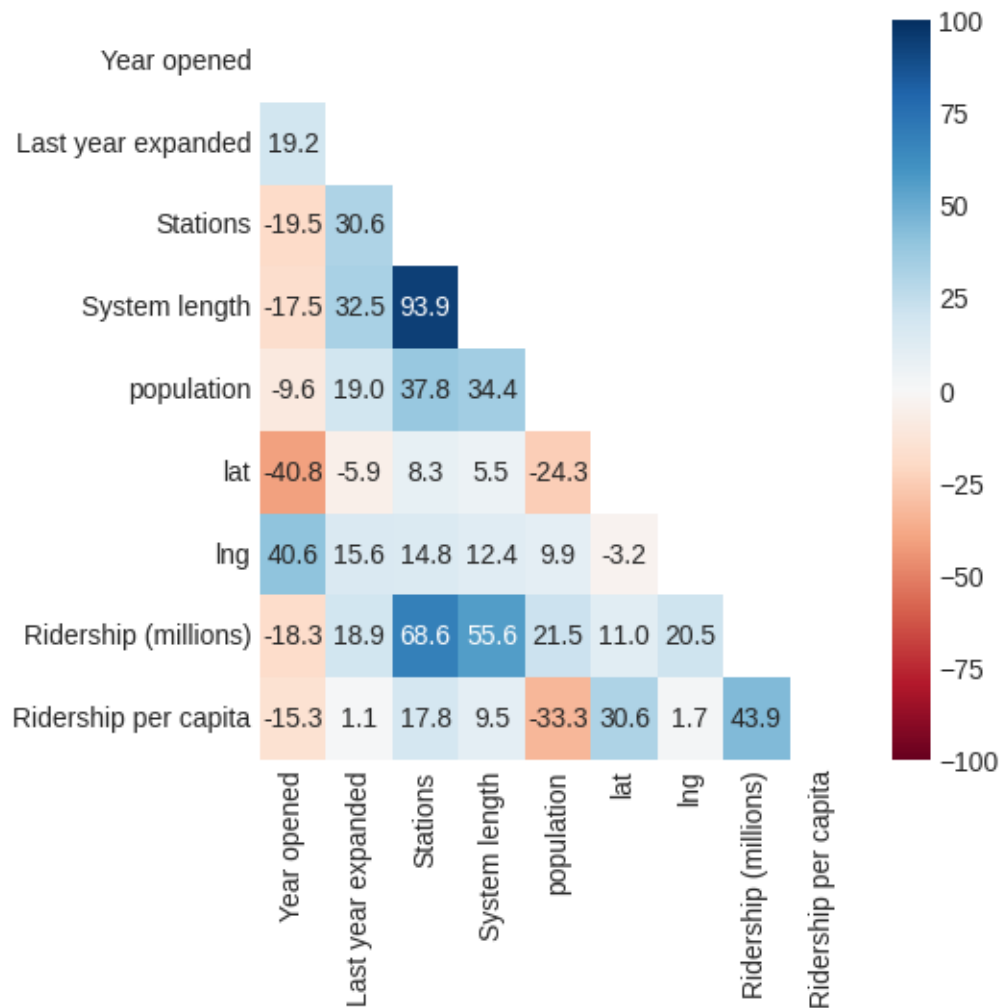
```
** Blancos **
```

```
Non zero cols:
```

```
Series([], dtype: int64)
```


```
Las correlaciones sin considerar outliers son:
```

```
[88]: corr_no_outliers = plot_corr(df_norm_no_outliers, cols_num);
```



Lo que reafirma la codependencia kilómetros-estaciones (94%):

```
[89]: g = sns.jointplot(data=df_norm_no_outliers, x='System length', y='Stations',
    ↪kind='reg')
g.fig.set_size_inches((3.5, 3));
```



La diferencia entre tabla de correlaciones con y sin outliers es moderada.

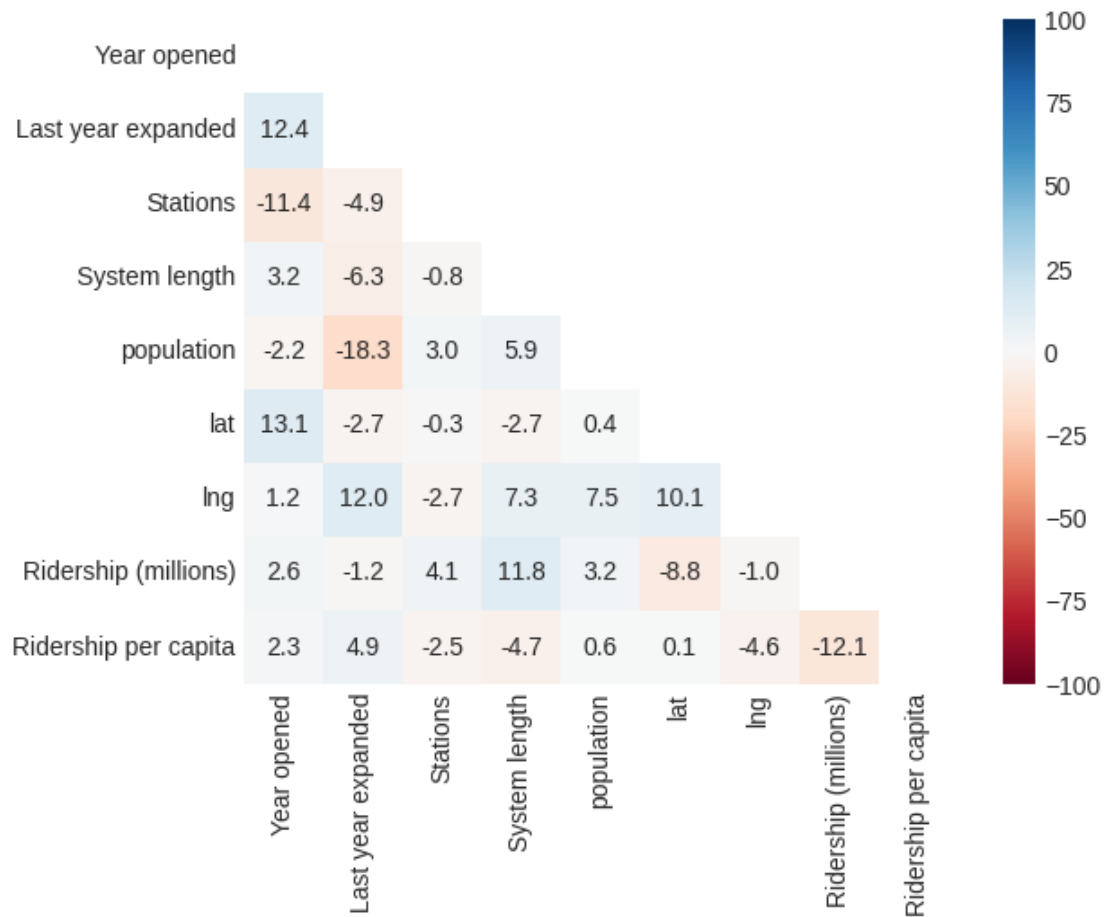
```
[90]: corr_diff = corr - corr_no_outliers
corr_diff.describe().loc[['mean', 'min', 'max']].applymap(lambda x: format(x, ".1f"))
```

```
[90]:
```

	Year opened	Last year expanded	Stations	System length	population	lat	\
mean	2.4	-0.5	-1.7	1.5	-0.0	1.0	
min	-11.4	-18.3	-11.4	-6.3	-18.3	-8.8	
max	13.1	12.4	4.1	11.8	7.5	13.1	

	lng	Ridership (millions)	Ridership per capita
mean	3.3	-0.2	-1.8
min	-4.6	-12.1	-12.1
max	12.0	11.8	4.9

```
[91]: sns.heatmap(corr_diff, annot=True, mask=np.triu(corr), fmt=".1f", cmap='RdBu',
                vmin=-100, vmax=100);
```



4.3 Regresión lineal

Para aplicar el modelo de regresión lineal se ha de trabajar con datasets sin NaN, pero los outliers pasaron a ser NaN, y con que fuese outlier en un campo ya he de eliminar la tupla entera:

```
[92]: df_norm_no_outliers_no_nan = df_norm_no_outliers.dropna()
```

```
[93]: from sklearn.linear_model import LinearRegression

model = LinearRegression(fit_intercept=True)

y_col = 'System length'
X = df_norm_no_outliers_no_nan.drop(y_col, axis=1).values
y = np.array(df_norm_no_outliers_no_nan[y_col])[:, np.newaxis]

X.shape, y.shape
```

```
[93]: ((112, 8), (112, 1))
```

Aplico el modelo

```
[94]: model.fit(X, y)
```

```
[94]: LinearRegression()
```

Este modelo indica que la aproximación \hat{y} al campo **System length** es tal que:

$$\hat{y} = -0.07 + 0.067 \cdot x_{\text{Last year expanded}} + 0.01 \cdot x_{\text{Year opened}} + 0.79 \cdot x_{\text{lng}} + 0.003 \cdot x_{\text{lat}} + 0.02 \cdot x_{\text{Ridership (millions)}} - 0.01 \cdot x_{\text{Stations}}$$

```
[95]: pd.DataFrame(data=model.coef_, columns=list(set(cols_num).
↳difference(set([y_col]))))
```

```
[95]:
```

	lat	Year opened	Stations	population	lng \
0	0.067262	0.010513	0.794723	0.002668	0.020939

	Ridership (millions)	Last year expanded	Ridership per capita
0	-0.01282	0.007279	-0.081639

```
[96]: model.intercept_
```

```
[96]: array([-0.07242396])
```

Con un coefficient de determinacion (R^2) de solo 93.5%, es decir, la regresión lineal es muy pobre para representar este campo, ó puede que dependa de otras varaibles no recogidas en nuestro dataset como PIB anual del país, inversión en infraestructuras per cápita, etc.

```
[97]: import statsmodels.api as sm
results = sm.OLS(y, X).fit()
print(results.summary(), '\n')
# individual results parameters can be accessed
print('Parameters: ', results.params)
print('R2: ', results.rsquared)
```

OLS Regression Results

```
=====
=====
Dep. Variable:                y    R-squared (uncentered):
0.935
Model:                    OLS    Adj. R-squared (uncentered):
0.930
Method:                Least Squares    F-statistic:
187.1
Date:                Fri, 16 Jun 2023    Prob (F-statistic):
3.96e-58
Time:                21:20:36    Log-Likelihood:
227.91
No. Observations:                112    AIC:
```

```

-439.8
Df Residuals:          104    BIC:
-418.1
Df Model:              8
Covariance Type:      nonrobust
=====
              coef    std err          t      P>|t|      [0.025    0.975]
-----
x1             0.0553     0.022     2.495     0.014     0.011     0.099
x2            -0.0482     0.029    -1.656     0.101    -0.106     0.010
x3             0.8075     0.044    18.402     0.000     0.720     0.895
x4            -0.0063     0.029    -0.221     0.825    -0.063     0.051
x5             0.0076     0.021     0.357     0.722    -0.035     0.050
x6            -0.0088     0.013    -0.668     0.505    -0.035     0.017
x7             0.0053     0.035     0.153     0.879    -0.063     0.074
x8            -0.0888     0.039    -2.293     0.024    -0.166    -0.012
=====
Omnibus:          64.171   Durbin-Watson:          2.037
Prob(Omnibus):    0.000   Jarque-Bera (JB):        344.509
Skew:            1.863   Prob(JB):                1.55e-75
Kurtosis:        10.742   Cond. No.                29.3
=====

```

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

Parameters: [ 0.055267 -0.0481955  0.80750771 -0.00634602  0.00760379
-0.00881157
 0.00527578 -0.08876115]
R2: 0.9350200760687406

```

4.4 Estudios diversos


Ciudades con mas de un sistema de metro:

```

[98]: unique_cities = df.groupby(['City', 'country']).size().reset_index(name='Count')
mask_cities = unique_cities['Count'] > 1

ax = sns.barplot(data=unique_cities[mask_cities], x='City', y='Count',
hue='country')
ax.legend(loc='center left', bbox_to_anchor=(1.0, 0.5))
plt.xticks(rotation=25);

```



Agrupamos datos de metros de una misma ciudad:

```
[99]: # total number of stations and system length for each city-country pair
city_country_totals = df.groupby(['City', 'country']).agg({'Stations': 'sum',
                                                           'System length': 'sum',
                                                           'Year opened': 'min',
                                                           'Last year expanded': 'max',
                                                           'lat': 'mean',
                                                           'lng': 'mean',
                                                           'Ridership (millions)': 'sum',
                                                           'Ridership per capita': 'sum'}).reset_index()
```

Por ejemplo Nueva York tiene 3 redes de metro, con un total de 458 estaciones:

```
[100]: df[df['City']=="New York City"]
```

```
[100]:
```

	City	Name	Year opened	Last year expanded	\
194	New York City	New York City Subway	1904	2017.0	
195	New York City	Staten Island Railway	1925	2017.0	
196	New York City	PATH	1908	1937.0	

	Stations	System length	country	admin_name	population	lat	\
194	424	399.0	United States	New York	18972871.0	40.6943	
195	21	22.5	United States	New York	18972871.0	40.6943	
196	13	22.2	United States	New York	18972871.0	40.6943	

	lng	Ridership (millions)	Ridership Year	Ridership per capita
194	-73.9249	NaN	NaN	NaN
195	-73.9249	3.8	2022	0.200286
196	-73.9249	45.5	2022	2.398161

```
[101]: city_country_totals[city_country_totals['City']=="New York City"]
```

```
[101]:
```

	City	country	Stations	System length	Year opened	\
120	New York City	United States	458	443.7	1904	

	Last year expanded	lat	lng	Ridership (millions)	\
120	2017.0	40.6943	-73.9249	49.3	

	Ridership per capita
120	2.598447


Redes de metro: cuantas se crean y cuando han sido por último expandidas.

```
[102]: fig, ax = plt.subplots(figsize=(20, 8))
cols = ['Year opened', 'Last year expanded']
for col in cols:
    df_aux = city_country_totals[col].value_counts().reset_index(name='Count')

    l = sns.lineplot(data=df_aux, x='index', y='Count', markers=True,
↳dashes=False, ax=ax, label=col)
    l.set(yscale="log")

plt.xticks(rotation=90)
plt.grid(visible=True, which='both', color='black', linewidth=0.075)

# Customize y-ticks
ax.set_yscale("log")
ax.yaxis.set_major_locator(ticker.LogLocator(base=10, subs=range(1,9)))
ax.yaxis.set_major_formatter(ticker.ScalarFormatter())
```

Cantidad de redes de metro que se han abierto cada año (`Year opened`), pero solo de los 5 países con mas redes. Similar con `Last year expanded`.

```
[103]: def plot_year(df=city_country_totals, col='Year opened', top_n=5):

    # occurrences of each country
    country_counts = df['country'].value_counts()

    # top 5 countries
    top_countries = country_counts.head(top_n).index



    # filter
    filtered_df = df[df['country'].isin(top_countries)]

    # group the filtered DataFrame
    grouped = filtered_df.groupby([col, 'country']).size().unstack()

    # plot
    ax = grouped.plot(kind='bar', stacked=True)
    ax.yaxis.set_major_locator(ticker.MaxNLocator(integer=True))
    ax.set_xticklabels([int(float(label.get_text()))
                        if float(label.get_text()).is_integer()
                        else label.get_text()
                        for label in ax.get_xticklabels()])

    plt.xlabel(col)
    plt.ylabel('Estaciones de metro')
    plt.legend(title='Country', bbox_to_anchor=(1.05, 1), loc='best')
    plt.plot()

[plot_year(col=k) for k in ['Year opened', 'Last year expanded']];
```



Países con mas estaciones:

```
[104]: def df_count_values(df=city_country_totals, col='Stations', top_n=10,
↪mean=False):
    # stations for each country


    if not mean:
        df = df.groupby('country')[col].sum().reset_index()
    else:
        df = df.groupby('country')[col].mean().reset_index()

    # top 10 countries with the most stations
    top_countries = df.nlargest(top_n, col)

    ax = sns.barplot(data=top_countries, x='country', y=col)
    for container in ax.containers:
        ax.bar_label(container, fmt='%.1f');
    plt.xticks(rotation=25)

    return top_countries;
```


```
[105]: top_countries = df_count_values();
```



En eje de abscisas situamos a la izquierda los países con mas estaciones, y dentro de cada país también ordenamos sus ciudades según el número de estaciones descendentes.



```
[106]: def df_plot_each_city(df=city_country_totals, col='Stations',  
    ↪top_countries=top_countries):  
  
    # Filter the original DataFrame for the selected countries  
    filtered_df = df[df['country'].isin(top_countries['country'])]  
  
    # sort cities in each country  
    filtered_df = filtered_df.sort_values(by=['country',col], ascending=[True,  
    ↪False])  
  
    # sort countries not unsorting its cities  
    boolVar = True  
    for k in top_countries['country']:  
        df_k = filtered_df[filtered_df['country'] == k]  
        if boolVar:  
            df_aux = df_k.copy()  
            boolVar = False  
        else:  
            df_aux = pd.concat([df_aux,df_k])  
  
    # bar plot  
    plt.figure(figsize=(20, 8))  
    ax = sns.barplot(data=df_aux, x='City', y=col, hue='country')  
  
    plt.xlabel('City')  
    plt.ylabel(col)  
    plt.title(f'Top 10 Countries with the Most {col}')  
    plt.legend(title='Country')  
    ax.legend(loc='center left', bbox_to_anchor=(0.1, -0.3),fancybox=True,  
    ↪shadow=True, ncol=4)  
    plt.xticks(rotation=90);
```

```
[107]: df_plot_each_city()
```




Países con mas kilómetros. Francia es la sexta con mas estaciones, pero la décima en kilómetros:

```
[108]: col="System length"
top_countries = df_count_values(col=col);
df_plot_each_city(col=col, top_countries=top_countries);
```



Países con mas viajeros (en el año recogido en la base de datos):

```
[109]: col="Ridership (millions)"
top_countries = df_count_values(col=col);
df_plot_each_city(col=col, top_countries=top_countries);
```




```


# Iterate over countries and plot stacked bar plots
for i, country in enumerate(top_countries['country']):
    data = df[df['country'] == country][[col, 'Ridership (millions)']]
    data[f'{col}_range'] = pd.cut(data[col], latlng_ranges) # Bin the
    ↪ latitude values into ranges
    grouped_data = data.groupby(f'{col}_range')['Ridership (millions)'].
    ↪ sum() # Group by lat_range and sum the ridership
    grouped_data = pd.DataFrame(grouped_data)


    sns.barplot(data=grouped_data, x=grouped_data.index, y='Ridership
    ↪ (millions)',
                ax=ax, bottom=stacked_heights, label=country,
    ↪ color=colors[i])
    stacked_heights += grouped_data['Ridership (millions)'].values

# Customize plot
plt.xlabel(f'{col} Range')
plt.ylabel('Ridership (millions)')
plt.title(f'Passengers by {col} Range and Country')
plt.legend(title='Países con más viajeros', bbox_to_anchor=(1.05, 1),
    ↪ loc='upper left')
plt.xticks(rotation=90)
plt.show()

[plot_lat_lng(col=k) for k in ['lat', 'lng']];

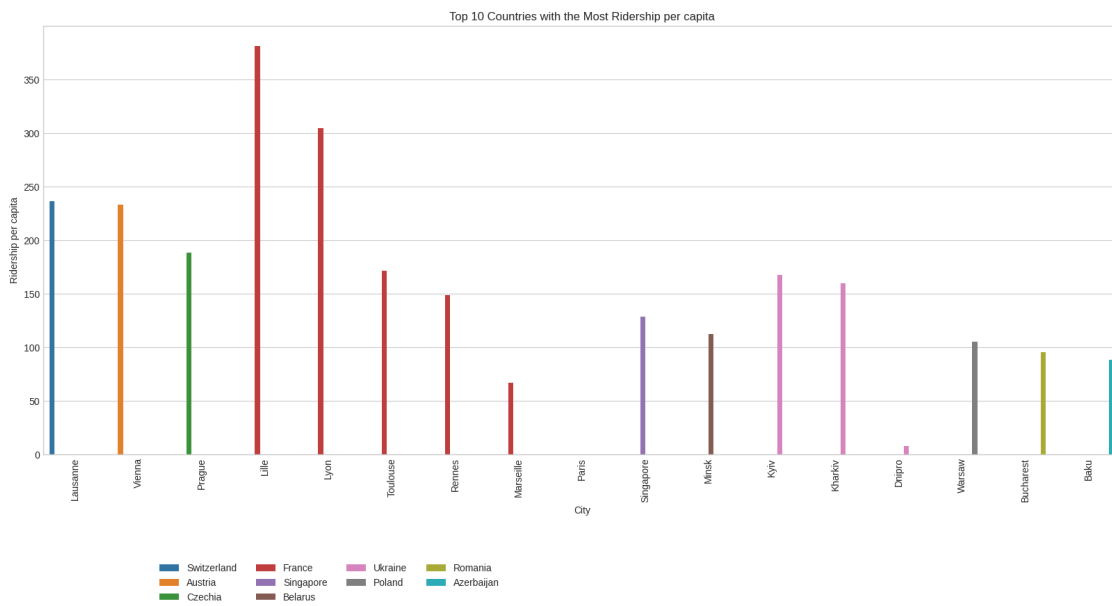

```





Países con mas viajeros per capita (en el año recogido en la base de datos).


```
[111]: col="Ridership per capita"
top_countries = df_count_values(col=col, mean=True);
df_plot_each_city(col=col, top_countries=top_countries);
```



Ciudades con mas estaciones:


```
[112]: def plot_top_cities(df=df, col='Stations'):  
        top_cities = df.groupby(['City', 'country'])[col].sum().nlargest(20).  
        ↪reset_index(name=col)  
        ax = sns.barplot(data=top_cities, x='City', y=col, hue='country')  
        ax.legend(loc='center left', bbox_to_anchor=(1.0, 0.5))  
        plt.xticks(rotation=90);
```

```
[113]: plot_top_cities()
```




Ciudades con mas kilómetros:

```
[114]: plot_top_cities(col="System length")
```




Ciudades con mas viajeros por año (en el año recogido en la base de datos):

```
[115]: plot_top_cities(col="Ridership (millions)")
```



Ciudades con mas viajeros per capita (en el año recogido en la base de datos).

```
[116]: plot_top_cities(col="Ridership per capita")
```



5 Resolución del problema y conclusiones.

La distribución de campos numéricos (violinplots) resumen muy visualmente la base de datos: no muchos datos (solo 205 registros) pero muy variados.

Los campos han sido preprocesados correctamente, restando solo 16 registros de viajeros no indicados. Ya que los NaN de la variable `Last year expanded` simplemente indican que hay 22 metros sin ampliaciones.

Los heatmaps de correlaciones con ó sin outliers indican una alta falta de correlación excepto entre estas 3 parejas de campos: `System length-Stations`, `System length-Ridership (millions)` y `Stations-Ridership (millions)`. Precisamente por tener esta ventana se ha elegido la variable `System length` para su estudio de predicción como regresión lineal en función del resto de variables numéricas. El moderado coef. de determinación del 93% no es suficientemente alto como para considerar el modelo satisfactorio. Si se deseara predecir esa variable se recomienda: - recurrir a modelos mas complejos como árboles de decisión dónde también participen variables categóricas como el país - aumentar el número de registro buscando en internet mas lejos del dataset de kaggle si hay mas metros a día de hoy - recurrir a otras bases de datos con datos relativos a la ciudad/país como PIB anual del país, inversión en infraestructuras per cápita, etc. ó al metro como coste de mantenimineto, personal operativo, grado de satisfacción de los viajeros,...

Por otro lado, los múltiples gráficos de barras posteriores a la regresión sí nos han aportado muy

valiosa información: - China está acelerando la construcción de líneas de metro desde el principio de siglo (la India potente también pero desde 2011), es donde mas se usa en términos absolutos, y desde 2019 es donde mas líneas se amplían - los países con mas estaciones, kilómetros de vías y viajeros siguen un patrón común de sus redes de metro, donde una megaurbe (Nueva York, Seul, Tokyo, Dehli,...) puntua claramente en cada uno de estos aspectos al resto de urbes de su país. Este sorpaso es mínimo el doble, llegando en algunas situaciones a ser triple. Después, el resto de metros del país (ciudades 3º, 4º, etc.) reducen sus cuantías de los 3 citados campos ya de una manera mas progresiva. Ésto no sucede en China (lider absoluto de estos 3 parámetros) donde Shangai, Pekín, Shenzhen ó Chengdu compiten mas a la par por la supremacía de este servicio - en cuanto al uso per cápita (pasajeros anuales entre población urbana) destacan países fuera de top absolutos como Suiza (con Lausanne), Austria (con Viena) y Rep. Checa (con Praga) - si analizamos las líneas sin prefiltrar que tengan que pertenecer a la cima en el aspecto en estudio, se observa que: - de las 20 con mas estaciones solo 5 son europeas: 5º Londres, 6º París, 8º Moscú, 14º Madrid y 19º Berlín - de las 20 mas largas solo 3 son europeas: 9º Moscú, 13º Londrés y 17º Madrid - de las 20 mas usadas solo 5 son europeas: 11º Kiev, 12º Berlín, 14º Viena, 15º Madrid, 16º Milán (aunque es un dato a revisar, ya que los datos pueden ser muy antiguos, ó por ejemplo no se han registrado los viajeros de Pekín) - de las 20 mas usadas per cápita 16 son europeas, con especial asiduididad se suben al metros los ciudadanos de Lille (Francia). Esta gráfica ha de tener las mismas consideraciones que la anterior.

Conclusiones: - el lejano oriente (y la India-Irán) ya es lider en transporte suburbano y su tendencia es fuertemente creciente - si extrapolamos datos del metro: en Estados Unidos y Europa suele haber unas pocas ciudades por país donde se concentra el desarrollo, mientras en los dragones asiáticos no es tan pronunciado el contraste - Europa lidera las ciudades donde cada habitante mas usa el metro

6 Exportación del código y de los datos producidos

Guardamos el DataFrame mergeado y limpiado (pero no normalizado) en un archivo CSV.

```
[117]: file_output = "../dataset/df_metro_cities.csv"
df.to_csv(file_output, index=False)
```

Lo cargamos para comprobar que se guardó satisfactoriamente.

```
[118]: pd.read_csv(file_output).tail(3)
```

```
[118]:
```

	City	Name	Year opened	Last year expanded	Stations	\
202	Tashkent	Tashkent Metro	1977	2023.0	48	
203	Caracas	Caracas Metro	1983	2015.0	49	
204	Hanoi	Hanoi Metro	2021	NaN	12	

	System length	country	admin_name	population	lat	\
202	59.1	Uzbekistan	Toshkent	2571668.0	41.3111	
203	67.2	Venezuela	Distrito Capital	2245744.0	10.4806	
204	13.1	Vietnam	Hà Nội	8246600.0	21.0283	

```
lng Ridership (millions) Ridership Year Ridership per capita
```


202	69.2797	136.7	2022.0	53.156162
203	-66.9036	NaN	NaN	NaN
204	105.8542	NaN	NaN	NaN

[]: